

Optimization-AI

Plain-English in → provably-optimal decision out. An LLM front-end over a real operations-research solver.

What it is, in one sentence

You describe a logistics or planning problem in ordinary language; the system figures out *what kind* of optimization problem it is, pulls the numbers out of the prose, runs an exact mathematical solver, and explains the answer back in plain English — with no spreadsheet, no modelling language, and no solver expertise required from the user.

Why this is worth something

Operations Research already *solves* these problems optimally — the hard part has always been the human bottleneck in between: someone has to translate a messy business question into a precise mathematical model. That step needs a specialist and takes hours. This system collapses that step to a sentence, while keeping the answer **exact** (it is a real solver, not the LLM guessing numbers).

What it can do today — two live examples

Example 1 Classic transportation (linear program)

"Two factories produce widgets: Seattle has capacity 350 and San Diego has capacity 600. Ship to New York (325), Chicago (300), Topeka (275). Distances in thousand miles ... freight is \$90 per case per thousand miles. Minimize total shipping cost."

Result: optimal shipping plan, total cost **\$153,675** — matches the textbook optimum exactly.

Example 2 Fixed-charge logistics (mixed-integer program)

"Ship from 2 plants to 3 customers... on top of per-unit cost, opening any route costs a fixed \$500 regardless of volume. Minimize total cost including the fixed charges."

Result: optimal cost **1660** — the solver decides it is cheaper to consolidate onto **three** routes

($P1 \rightarrow M1$, $P1 \rightarrow M2$, $P2 \rightarrow M3$) and leave the rest closed. This is a genuinely hard combinatorial decision (which routes to open), not just arithmetic.

How it works

Every stage is a discrete, inspectable step — not one opaque prompt:

natural language

- > Intent router (new problem / follow-up / chit-chat)
- > Classifier (transportation? scheduling? – majority vote, schema-bound)
- > Specialist extractor (prose -> typed parameter table)
- > Feasibility check (3 layers: structural -> domain rules -> LP relaxation)
- > Solver (Pyomo + HiGHS – exact LP / MIP)
- > Explanation (plain-language summary, numbers grounded in the solution)

The feasibility gate matters: if the LLM hallucinates a number, the structured checks catch the inconsistency *before* the solver runs, so a bad extraction fails loudly instead of producing a confident wrong answer.

What is solid vs. honest limitations

Solid today

- Transportation (LP) & fixed-charge (MIP) — exact, verified against textbook optima
- Single-stage scheduling
- Deterministic numeric output (same answer every run, any machine)
- Runs fully local & private (Ollama / qwen3:8b) — no data leaves the box
- Round-trip eval harness: 3/3 pass, recall 1.0
- Classifier choice validated against ML, RAG & taxonomy baselines — three approaches built and benchmarked, simplest won on evidence (§8)

Honest limitations

- First answer is slow on CPU (~2–3 min of local LLM inference) — the next roadmap item makes the wait visible/streamed
- Two problem families so far; more to add
- Heuristic warm-start is an honest negative on current formulations (kept for the interactive UX + LP bound, not as a speedup — §9)
- Explanation wording varies slightly run-to-run (the *numbers* do not)

The high-level picture ends here. The rest of this document is for anyone who wants to dive into *how* it

actually works, the modelling assumptions, and what is next.

Part II — Technical deep dive

For the reader who wants the internals.

1. How a request is orchestrated

A single entry point (`agent/core.py` → `solve_natural_language`) runs a fixed, inspectable pipeline. Every stage is a discrete step with its own module, not one opaque megaprompt:

1. **Intent router** — classifies the message as *smalltalk* / *help* / *follow-up* / *optimization*. The first three short-circuit; only a genuine problem enters the solve pipeline.
2. **Problem classifier** — picks the problem family and the concrete `solver_id` using an *N-vote majority* over a JSON-schema-bound prompt. Low confidence or "UNKNOWN" stops here with a clear message rather than guessing.
3. **Specialist extractor** — a per-family extractor (`TransportationSpecialist` / `SchedulingSpecialist`) turns prose into a typed parameter dict, mirroring an example schema and omitting optional fields (e.g. `fixed_cost`) when not mentioned.
4. **Validation + 3-layer feasibility** — structural checks, then problem-specific necessary conditions (e.g. total supply $\geq$ total demand), then an LP-relaxation solve. An infeasible instance returns plain-language reasons *and concrete repair suggestions*, and is remembered so the next message can fix it.
5. **Mode routing** — `exact`, `heuristic`, or `heuristic_then_ask` (see §3).
6. **Solve → explain** — Pyomo + HiGHS produces the numbers; the LLM narrates them. The solution is stored in the conversation context so follow-ups and what-ifs work against it.

2. Why a pipeline, not an agent framework

The orchestration above is a **deterministic decision tree written in Python**, not an LLM-driven agent loop, and not LangChain. This is a deliberate design choice, not an omission — so it is worth stating plainly what was declined and why.

The governing principle. The LLM is confined to the *edges* of the system as a constrained natural-language ↔ structured-data adapter: prose in → typed parameter dict; numeric solution out → prose. It never decides *what the system does next*. The control flow — intent → classify → extract → validate → feasibility → solve → explain — is fixed and known at design time, and the actual optimisation is performed by a provably correct solver (Pyomo + HiGHS), not by the model.

The hard, unreliable boundary (free text → numbers) is tamed with JSON-schema-bound prompts, Pydantic validation, and the 3-layer feasibility gate (§7) — not with an agent that retries until something sticks.

What an agent framework actually buys you is two things: (1) an *agent-executor loop* where the model chooses tools and chains them dynamically because the right sequence cannot be enumerated ahead of time, and (2) retrieval plumbing (loaders, splitters, vector stores, retrievers) for RAG. Neither applies here. The tool set is a small, fixed registry of solvers; the right one is picked by a deterministic classifier + registry lookup, not chosen by the LLM in a loop. And retrieval was built, benchmarked, and *rejected on evidence* (RAG 50% vs 70% no-RAG; see §8) — it is archived, not in the production path. A framework would add abstraction, indirection, and a harder-to-debug failure surface for capabilities the system does not use.

The precise trigger — when this decision would flip. The argument is *not* “LLM output is stochastic, therefore avoid agents”; token-level non-determinism is handled by schema + validation + the feasibility gate. An agent loop earns its complexity when the **LLM itself must do dynamic planning over an open-ended action space** — specifically: the tool space is large and the correct sequence depends on intermediate results in ways the developer cannot pre-wire; the plan must be discovered at runtime; multi-hop retrieval is in the loop; or the task is long-horizon and self-correcting over many steps. This domain has none of those properties: it is bounded, the optimal control flow is knowable in advance, and the heavy lifting is operations research, not LLM orchestration. Adding more problem families or follow-up types does *not* cross that line — it is one more registered solver and one more classifier label, still a decision tree. If the tool space became open-ended or the plan had to be discovered per request, *that* is the point at which an agent-executor (LangGraph-style) loop would be the right tool. Until then, deterministic, code-owned orchestration is the more reliable, inspectable, and debuggable choice — and inspectability is itself a feature for a system whose outputs are business decisions.

3. The two-call heuristic protocol

Local LLM inference costs ~2–3 min on CPU. Rather than make the user stare at a spinner, `heuristic_then_ask` mode splits the work:

- **Call 1:** a fast constructive heuristic returns a feasible answer immediately, together with an *LP-relaxation lower bound* → an honest “this is at most X% above optimal” gap. The job is persisted in an in-memory store (10-minute TTL).
- **The user decides:** `optimize` / `accept` / `use heuristic` — or just asks a free-text what-if, which is answered against the pending job without consuming it.

- **Call 2 (optimize):** the exact solver runs *warm-started* from the heuristic's solution → a provably-optimal answer, reached faster.

How the constructive heuristic works

Transport — Vogel's Approximation (VAM). For every unsatisfied row and column it computes a *penalty* = the difference between the two cheapest available arcs in that line (the regret of *not* using the cheapest one). It picks the line with the largest penalty, ships as much as possible through that line's lowest-cost cell, removes the now-exhausted supply or demand, and repeats. This regret-greedy rule lands far closer to the optimum than a naive least-cost fill, in $O(n^2)$ with no solver.

Scheduling — Longest-Processing-Time (LPT). Jobs are sorted by descending processing time and each is placed on the machine that is currently free earliest. Front-loading the long jobs is the classic $(4/3 - 1/3m)$ -approximation for makespan and yields the assignment + sequence + completion times used as the warm start.

The fixed-charge correction (a real subtlety). VAM optimises variable cost only — it is blind to the per-route fixed charge f_{ij} . If we reported only its variable cost, the heuristic would look *cheaper* than the LP bound and the gap would go nonsensically negative. So the reported cost is `vam.cost` **plus one f_{ij} for every route VAM actually opens**, making the heuristic-vs-bound comparison sound.

Where the “X% above optimal” number comes from

The same MIP is solved once with the binaries relaxed ($y_{ij} \in [0, 1]$) — an LP, fast, and a *true lower bound* on the integer optimum. The honest gap is then $(\text{heuristic_cost} - \text{lp_bound}) / \text{heuristic_cost}$. If it is ≈ 0 the heuristic already matches the bound → reported as *provably optimal* and the exact solve is skipped entirely. Scheduling has no LP bound wired yet, so it honestly reports the makespan with no gap claim rather than a fabricated one.

How the warm start is actually injected

Call 1 persists a `JobRecord` = extracted params + the heuristic flows/assignment + the LP bound (10-min TTL). On `optimize` the exact solver is called with `warm_start=record.heuristic_flows`: each Pyomo decision variable is seeded via `var.value=...` *before* solve, and HiGHS through Pyomo's APPSI interface consumes that as a MIP start — it begins branch&bound from a known-feasible incumbent instead of from scratch. `accept / use heuristic` short-circuit with no second solve at all; a what-if question is evaluated against the still-pending job without consuming it.

4. How the optimization is actually solved

Models are built with **Pyomo** and solved by **HiGHS** through its `highspy` wheel via Pyomo's APPSI interface — no external solver binary to install (this is also why Windows setup is now trivial).

Transportation — linear program

```
min  Σij cij · xij
s.t.  Σj xij ≤ capacityi      (supply)
      Σi xij ≥ demandj      (demand)
      xij ≥ 0      (optional per-arc capacity xij ≤ uij)
```

Fixed-charge logistics — mixed-integer program

```
min  Σij cij · xij + Σij fij · yij
s.t.  transportation constraints, plus
      xij ≤ Mij · yij ,    yij ∈ {0,1}
      Mij = min(capacityi, demandj)    (tight big-M)
```

The binary `yij` = "is route (i,j) open?" turns this into a real combinatorial decision. Supplying a `fixed_cost` automatically switches the model to a MIP, enables the LP-relaxation bound, and trips the warm-start gate. Single-stage scheduling is modelled analogously (assignment + sequencing, makespan objective; LPT heuristic for the warm answer).

Determinism: HiGHS returns the exact optimum and classification/extraction run at temperature 0, so the objective and flows are byte-identical on any machine; only the narration wording can vary.

5. The LLM layer

An abstract `LLMClient` hides the provider. The default is local Ollama with `qwen3:8b` across all stages (one warm model); a Groq backend is opt-in. The model was chosen by an evidence-driven sweep (`evals/model_sweep.py`): `qwen3:8b` matched the larger `qwen3:14b` on every accuracy metric across both domains at roughly **twice the speed**, so it became the default. Per-stage temperature: classification 0 (schema-bound, voted), extraction 0, explanation 0.1–0.3 (wording only).

The key design rule: **the LLM never computes the optimum**. It only translates prose ↔ schema and narrates the solver's output. The 3-layer feasibility gate is the guardrail — a hallucinated parameter fails loudly *before* the solver runs, instead of producing a confident wrong number.

6. Follow-ups & what-if

Follow-ups are classified into *sensitivity* / *what-if* / *resolve* / *pareto*. A what-if re-parses just the changed parameter, re-checks feasibility, re-solves, and reports the cost delta and flow changes against the baseline. An infeasible scenario is treated as a useful answer ("that can't work, here's why, here's how to fix it"), not an error — and it works even while a heuristic job is still pending.

7. The 3-layer feasibility gate

Between extraction and the solver sits a deliberately ordered gate: **cheapest, most certain checks first; the expensive solver call last and only if needed**. Each layer can reject with plain-language reasons *and concrete repair suggestions*; the rejected instance is remembered so the next message can fix it.

Layer	What it does	Cost	Example catch
0 — Structural / sanity (problem-agnostic, pure Python)	Empty index sets, None/empty keys, dimensional consistency (cost matrix = sources × sinks), finite & non-negative values	microseconds	capacity = -50; cost matrix of the wrong shape
1 — Problem-specific necessary conditions (one <code>FeasibilityPlugin</code> per domain)	Each domain registers <i>one</i> plugin bundling its checker and its repair-suggestion generator, so a domain cannot be half-wired. Transport: total supply ≥ demand, per-sink capacity. Scheduling: an order's fastest eligible processing time ≤ its due date; no order without an eligible unit	microseconds	demand 900 > supply 700 → 200 shortfall; "OrderC needs 5 h, due hour 4"
2 — Solver-based (LP relaxation, the backstop)	Builds the real Pyomo model via the solver's <code>build_model</code> , relaxes binary/integer → continuous, solves <i>only for feasibility</i> . Conclusive for both transportation and scheduling, and automatic for any future solver exposing	one LP solve	interaction / per-arc-capacity or joint-deadline patterns Layers 0–1 cannot see

`build_model`

Layer 2 returns FEASIBLE / INFEASIBLE / UNKNOWN. The gate is **fail-closed**: an inconclusive layer yields `UNKNOWN` — it never asserts "feasible from ignorance" — and defers to the hardened, fail-soft solver as the real backstop. The report carries `status`, `reasons`, `layer_passed` and actionable suggestions. The gate is regression-pinned by a curated, machine-checkable infeasible corpus (transportation + scheduling, one case per layer) that asserts each instance is rejected at exactly its intended layer.

Why this matters. This is the guardrail against LLM hallucination. If the extractor invents or misreads a parameter, the structured checks catch the inconsistency *before* the solver runs — so the system fails loudly with a fixable explanation instead of producing a confident wrong number. Fail fast, fail cheap, fail actionable.

8. Classification — three approaches evaluated

Production uses an **LLM 3-vote majority** classifier. It was not the first or only thing tried: two alternatives were fully built and benchmarked first, and a third was engineered as a deterministic fallback.

Approach	Accuracy on real OR problems	Speed	Status
LLM, 3-vote majority	70%	3–5 s	shipped
Random-Forest + TF-IDF	44% (95% train / ~75% held-out)	<1 ms	rejected
LLM + RAG retrieval	50%	~35 s, timeouts	rejected
Taxonomy + weak-supervision hybrid	built (deterministic, auditable)	<1 ms	dormant fallback

- **ML classifier:** Random Forest over a curated **523-instance** dataset (OR-Library benchmarks + Chain-of-Experts/LPWP + synthetic; 24 subtypes, 11 families). It memorised the training distribution (95% train) but generalised poorly to real OR-Library problems (44% vs the LLM's 70%). Ensemble (LLM+ML) and two-stage (ML suggests → LLM verifies) *also* failed to beat the LLM alone — the errors were correlated, not complementary.

- **RAG:** LangChain + Chroma over 8 OR textbooks/papers (**20,399 chunks, 5,008 pages**, all-MiniLM-L6-v2 embeddings). It made classification *worse* (50% vs 70%) and caused extraction timeouts (2/10 problems >60 s): generic textbook definitions conflicted with the specific problem and added noise, not signal. RAG helped factual recall, not structural classification.
- **Taxonomy + weak supervision** (`or_classify/`): a versioned, hierarchical **9-family OR taxonomy** — each family carrying definitions, subtypes, examples and explicit *near-misses* — with **7 Snorkel-style labeling functions** and a hybrid feature pipeline (TF-IDF + LF votes + engineered features). Designed as a deterministic, auditable, schema-grounded alternative; built and retained as a fallback, not in the production path.
- **Takeaway:** the simplest option — an LLM 3-vote majority — was both the most accurate and the cheapest to operate. The elaborate approaches are preserved as evidence and a deterministic fallback. A documented "simple baseline is hard to beat" outcome.

9. Warm-start — an honest negative result

The plan assumed a heuristic warm-start would accelerate every solve. It does not. Warm-start helps only small pure-LP transport; beyond $\sim 10 \times 20$ it is *slower* than a cold solve, and on scheduling the cold and warm runs finish within milliseconds of each other — because the bottleneck is the MIP optimality-proof gap, not finding a feasible point. Consequence: warm-start is now **gated to the fixed-charge MIP** and off elsewhere. Separately, HiGHS was chosen over GLPK by a gate experiment (cold solve 79 ms on the reference instance, honours `time_limit / mip_rel_gap`, installs as a wheel — no external binary).

Assumptions & scope

Being explicit about what the current system assumes:

- The LLM extracts and explains; **all arithmetic is the solver's**.
- Transportation is single-commodity, single-period; costs are linear in flow; capacities, demands and costs are point estimates (no uncertainty/stochastics).
- A fixed charge, when present, is incurred once per opened route regardless of volume; $M_{ij} = \min(\text{capacity}_i, \text{demand}_j)$ is taken as a valid tight bound.
- The fast feasibility checks assume standard supply/demand structure; the LP-relaxation layer is the backstop for anything they miss.
- Scheduling is single-stage.
- Reproducibility assumes the same model (`qwen3:8b`) and solver (HiGHS); explanation

prose may differ run to run.

- Single-process, in-memory job store with a 10-minute TTL — not yet horizontally scaled.
- Demo posture: trusted input, no authentication / rate-limiting / multi-tenancy.

Future work

- **Make the wait visible** — stream the classify/extract/solve pipeline live so the LLM latency reads as progress, not dead air.
- **Follow-up chips & graceful fallback** — one-click what-if / sensitivity on the last solution.
- **Spreadsheet / API fast path** — upload an Excel sheet (or post structured params) and skip both LLM calls entirely; Excel export of results.
- **Fixed-charge-aware heuristic** — the reported heuristic cost is now correct, but the warm start (VAM) is still fixed-charge-blind; a charge-aware constructive heuristic would tighten the gap.
- **More problem families** — multi-period / multi-commodity flow, richer scheduling, beyond the current two.
- **Evaluation depth** — extend the round-trip eval harness (Phase 3) and add adversarial extraction tests.
- **Productionization** (if it graduates from demo) — persistent/scalable job store, auth, rate limiting.